

6.824 2021 Lecture 17: Scaling Memcache at Facebook

Scaling Memcache at Facebook, by Nishtala et al, NSDI 2013

why are we reading this paper?

it's an experience paper, not about new ideas/techniques

three ways to read it:

- cautionary tale about not taking consistency seriously from the start
- impressive story of super high capacity from mostly-off-the-shelf s/w
- fundamental struggle between performance and consistency
- we can argue with their design, but not their success

how do web sites cope as they get more users?

a typical story of evolution over time:

1. single machine w/ web server + application + DB
 - DB provides persistent storage, crash recovery, transactions, SQL
 - application queries DB, formats HTML, &c
 - but: as load grows, application takes too much CPU time
2. many web FEs, one shared DB
 - an easy change, since web server + app already separate from storage
 - FEs are stateless, all sharing (and concurrency control) via DB
 - stateless -> any FE can serve any request, no harm from FE crash
 - but: as load grows, need more FEs, soon single DB server is bottleneck
3. many web FEs, data sharded over cluster of DBs
 - partition data by key over the DBs
 - app looks at key (e.g. user), chooses the right DB
 - good DB parallelism if no data is super-popular
 - painful -- cross-shard transactions and queries probably don't work
 - hard to partition too finely
 - but: DBs are slow, even for reads, why not cache read requests?
4. many web FEs, many caches for reads, many DBs for writes
 - cost-effective b/c read-heavy and memcached 10x faster than a DB
 - memcached just an in-memory hash table, very simple
 - complex b/c DB and memcacheds can get out of sync
 - fragile b/c cache misses can easily overload the DB
 - (next bottleneck will be DB writes -- hard to solve)

the big facebook infrastructure picture

lots of users, friend lists, status, posts, likes, photos

fresh/consistent data not critical

humans are tolerant

high load: billions of operations per second

that's 10,000x the throughput of one DB server

~100,000 simple queries/s for mysql

~1,000s transaction/s

~1000,000 get/puts/s for memcached

multiple data centers (at least west and east coast)

each data center -- "region":

"real" data sharded over MySQL DBs

memcached layer (mc)

web servers (clients of memcached)

each data center's DBs contain full replica

west coast is primary, others are secondary replicas via MySQL async log replication

how do FB apps use mc? Figure 1.

FB uses mc as a "look-aside" cache

real data is in the DB

cached value (if any) should be same as DB

read:

```
v = get(k) (computes hash(k) to choose mc server)
if v is nil {
  v = fetch from DB
  set(k, v)
}
```

```

write:
  v = new value
  send k,v to DB
  delete(k)
application determines relationship of mc to DB
mc doesn't know anything about DB

```

what does FB store in mc?
 paper does not say
 maybe userID -> name; userID -> friend list; postID -> text; URL -> likes
 data derived from DB queries

paper lessons:
 look-aside caching is trickier than it looks -- consistency
 paper is trying to integrate mutually-oblivious storage layers
 cache is critical:
 not really about reducing user-visible delay
 mostly about shielding DB from huge overload!
 human users can tolerate modest read staleness
 stale reads nevertheless potentially a big headache
 want to avoid unbounded staleness (e.g. missing a delete() entirely)
 want read-your-own-writes
 more caches -> more sources of staleness
 huge "fan-out" => parallel fetch, in-cast congestion

let's talk about performance first
 majority of paper is about avoiding stale data
 but staleness only arose from performance design

performance comes from parallelism due to many servers
 many active users, many web servers (clients)
 two basic parallel strategies for storage: partition vs replication

will partition or replication yield most mc throughput?
 partition: divide keys over mc servers
 replicate: divide clients over mc servers
 partition:
 + more memory-efficient (one copy of each k/v)
 + works well if no key is very popular
 - each web server must talk to many mc servers (overhead)
 replication:
 + good if a few keys are very popular
 + fewer TCP connections
 - less total data can be cached

performance and regions (Section 5)

[diagram: west, db primary shards, mc servers, clients | east, db secondary shards, ...
 feed from db primary to secondary]

Q: what is the point of regions -- multiple complete replicas?
 lower RTT to users (east coast, west coast)
 quick local reads, from local mc and DB
 (though writes are expensive: must be sent to primary)
 maybe hot replica for main site failure?

Q: why not partition users over regions?
 i.e. why not east-coast users' data in east-coast region, &c
 then no need to replicate: might cut hardware costs in half!
 but: social net -> not much locality
 might work well for e.g. e-mail

Q: why OK performance despite all writes forced to go to the primary region?
 writes are much rarer than reads
 perhaps 100ms to send write to primary, not so bad for human users

users do not wait for all effects of writes to finish
i.e. for all stale cached values to be deleted

performance within a region (Section 4)

[diagram: db shards, multiple clusters, each w/ mc's and clients]

multiple mc clusters *within* each region
cluster == complete set of mc cache servers
i.e. a replica, at least of cached data

why multiple clusters per region?

why not add more and more mc servers to a single cluster?

1. adding mc servers to cluster doesn't help single popular keys replicating (one copy per cluster) does help
2. more mcs in cluster -> each client req talks to more servers and more in-cast congestion at requesting web servers
client requests fetch 20 to 500 keys! over many mc servers
MUST request them in parallel (otherwise total latency too large)
so all replies come back at the same time
network switches, NIC run out of buffers
3. hard to build network for single big cluster
uniform client/server access
so cross-section b/w must be large -- expensive
two clusters -> 1/2 the cross-section b/w

but -- replicating is a waste of RAM for less-popular items
"regional pool" shared by all clusters
unpopular objects (no need for many copies)
the application s/w decides what to put in regional pool
frees RAM to replicate more popular objects

bringing up new mc cluster is a performance problem
new cluster has 0% hit rate
if clients use it, will generate big spike in DB load
if ordinarily 1% miss rate,
adding "cold" second cluster will causes misses for 50% of ops.
i.e. 50x spike in DB load!
thus the clients of new cluster first get() from existing cluster (4.3)
and set() into new cluster
basically lazy copy of existing cluster to new cluster
better 2x load on existing cluster than 30x load on DB

another overload problem: thundering herd
one client updates DB and delete()s a key
lots of clients get() but miss
they all fetch from DB
they all set()
not good: needless DB load
mc gives just the first missing client a "lease"
lease = permission to refresh from DB
mc tells others "try get() again in a few milliseconds"
effect: only one client reads the DB and does set()
others re-try get() later and hopefully hit

what if an mc server fails?

can't have DB servers handle the misses -- too much load
can't shift load to one other mc server -- too much
can't re-partition all data -- time consuming
Gutter -- pool of idle mc servers, clients only use after mc server fails
after a while, failed mc server will be replaced

The Question:

why don't clients send invalidates to Gutter servers?
my guess: would double delete() traffic

and send too many delete()s to small gutter pool
 since any key might be in the gutter pool

important practical networking problems:

- n² TCP connections is too much state
- thus UDP for client get()s
- UDP is not reliable or ordered
- thus TCP for client set()s
- and mcrouter to reduce n in n²
- single request per packet is not efficient (for TCP or UDP)
- per-packet overhead (interrupt &c) is too high
- thus mcrouter batches many requests into each packet

let's talk about consistency now

what is their consistency goal?

- writes go direct to primary DB, with transactions, so writes are consistent
- what about reads?
- reads do not always see the latest write
- e.g. since not guaranteed across clusters
- more like "not more than a few seconds stale"
- i.e. eventual
- *and* writers see their own writes (due to delete())
- read-your-own-writes is a big driving force

first, how are DB replicas kept consistent across regions?

- one region is primary
- primary DBs distribute log of updates to DBs in secondary regions
- secondary DBs apply
- secondary DBs are complete replicas (not caches)
- DB replication delay can be considerable (many seconds)

how do they keep mc content consistent w/ DB content?

1. DBs send invalidates (delete()s) to all mc servers that might cache
 this is McSqueal in Figure 6
2. writing client also invalidates mc in local cluster
 for read-your-own-writes

they ran into a number of DB-vs-mc consistency problems

- due to races when multiple clients read from DB and put() into mc
- or: there is not a single path along which updates flow in order

what were the races and fixes?

Race 1:

- k not in cache
- C1 get(k), misses
- C1 v1 = read k from DB
- C2 writes k = v2 in DB
- C2 delete(k)
- C1 set(k, v1)
- now mc has stale data, delete(k) has already happened
- will stay stale indefinitely, until k is next written
- solved with leases -- C1 gets a lease from mc, C2's delete() invalidates lease,
- so mc ignores C1's set
- key still missing, so next reader will refresh it from DB

Race 2:

- during cold cluster warm-up
- remember: on miss, clients try get() in warm cluster, copy to cold cluster
- k starts with value v1
- C1 updates k to v2 in DB
- C1 delete(k) -- in cold cluster
- C2 get(k), miss -- in cold cluster
- C2 v1 = get(k) from warm cluster, hits

C2 set(k, v1) into cold cluster
 now mc has stale v1, but delete() has already happened
 will stay stale indefinitely, until key is next written
 solved with two-second hold-off, just used on cold clusters
 after C1 delete(), cold mc ignores set()s for two seconds
 by then, delete() will (probably) propagate via DB to warm cluster

Race 3:

k starts with value v1
 C1 is in a secondary region
 C1 updates k=v2 in primary DB
 C1 delete(k) -- local region
 C1 get(k), miss
 C1 read local DB -- sees v1, not v2!
 later, v2 arrives from primary DB
 solved by "remote mark"
 C1 delete() marks key "remote"
 get() miss yields "remote"
 tells C1 to read from *primary* region
 "remote" cleared when new data arrives from primary region

Q: aren't all these problems caused by clients copying DB data to mc?
 why not instead have DB send new values to mc, so clients only read mc?
 then there would be no racing client updates &c, just ordered writes

A:

1. DB doesn't generally know how to compute values for mc
 generally client app code computes them from DB results,
 i.e. mc content is often not simply a literal DB record
2. would increase read-your-own writes delay
3. DB doesn't know what's cached, would end up sending lots
 of values for keys that aren't cached

PNUTS does take this alternate approach of primary-updates-all-copies

FB/mc lessons for storage system designers?

cache is vital for throughput survival, not just to reduce latency
 need flexible tools for controlling partition vs replication
 need better ideas for integrating storage layers with consistency

--- references

<http://cs.cmu.edu/~beckmann/publications/papers/2020.osdi.cachelib.pdf>